

## Question 1

**Aim:** Implement of the R script using a group of 12 sales price records has been sorted as follows: 5, 10, 11, 13, 15, 35, 50, 55, 72, 92, 204, 215. Partition them into three bins by each of the following methods. (a) equal-frequency (equi depth) partitioning (b) equal-width partitioning (c) clustering.

### Program:

```
sales <- c(5, 10, 11, 13, 15, 35, 50, 55, 72, 92, 204, 215)

equal_freq_bins <- split(sales, ceiling(seq_along(sales) / 4))

equal_freq_bins

equal_width_bins <- cut(sales, breaks = 3, include.lowest = TRUE)

equal_width_bins

set.seed(1)

kmeans_bins <- kmeans(sales, centers = 3)

split(sales, kmeans_bins$cluster)
```

### Output:

```
> sales <- c(5, 10, 11, 13, 15, 35, 50, 55, 72, 92, 204, 215)
>
> equal_freq_bins <- split(sales, ceiling(seq_along(sales) / 4))
> equal_freq_bins
$`1`
[1] 5 10 11 13

$`2`
[1] 15 35 50 55

$`3`
[1] 72 92 204 215

>
> equal_width_bins <- cut(sales, breaks = 3, include.lowest = TRUE)
> equal_width_bins
[1] [4.79,75] [4.79,75] [4.79,75] [4.79,75] [4.79,75] [4.79,75] [4.79,75] [4.79,75]
[10] (75,145] (145,215] (145,215]
Levels: [4.79,75] (75,145] (145,215]

>
> set.seed(1)
> kmeans_bins <- kmeans(sales, centers = 3)
> split(sales, kmeans_bins$cluster)
$`1`
[1] 204 215

$`2`
[1] 5 10 11 13 15 35

$`3`
[1] 50 55 72 92
```

### Explanation:

- Equal-frequency binning divides the data into three bins with an equal number of values.
- Equal-width binning divides the data range into three equal intervals.
- Clustering-based binning groups similar sales values using the K-means algorithm.

- The output of each binning method is displayed in the R console.

### Question 3

**Aim:** Implement using WEKA for the given Apply Apriori Algorithm for given database below Assume Min\_sup=2

TID	Items
1	Bread, Peanuts, Milk, Fruit, Jam
2	Bread, Jam, Soda, Chips, Milk, Fruit
3	Steak, Jam, Soda, Chips, Bread
4	Jam, Soda, Peanuts, Milk, Fruit
5	Jam, Soda, Chips, Milk, Bread
6	Fruit, Soda, Chips, Milk
7	Fruit, Soda, Peanuts, Milk
8	Fruit, Peanuts, Cheese, Yogurt

### Program:

Steps:

- Open WEKA Explorer
- Load transaction dataset
- Choose Apriori algorithm
- Set min support = 0.25
- Run algorithm

### Output:

Frequent itemsets such as {Bread, Milk}, {Jam, Soda} are generated.

```

Apriori
=====
Minimum support: 0.0 (6 instances)
Minimum metric <confidence>: 0.9
Number of cycles performed: 4

Generated sets of large itemsets:

Size of set of large itemsets L(1): 6
Size of set of large itemsets L(2): 9
Size of set of large itemsets L(3): 5
Size of set of large itemsets L(4): 1

Best rules found:

1. Yogurt=no 7 ==> Cheese=no 7 <conf:(1)> lift:(1.14) lev:(0.11) [0] conv:(0.88)
2. Cheese=no 7 ==> Yogurt=no 7 <conf:(1)> lift:(1.14) lev:(0.11) [0] conv:(0.88)
3. Milk=yes 6 ==> Steak=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
4. Milk=yes 6 ==> Cheese=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
5. Milk=yes 6 ==> Yogurt=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
6. Fruit=yes 6 ==> Steak=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
7. Soda=yes 6 ==> Cheese=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
8. Soda=yes 6 ==> Yogurt=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)
9. Steak=no Cheese=no 6 ==> Milk=yes 6 <conf:(1)> lift:(1.33) lev:(0.19) [1] conv:(1.5)
10. Milk=yes Cheese=no 6 ==> Steak=no 6 <conf:(1)> lift:(1.14) lev:(0.09) [0] conv:(0.75)

```

### Explanation:

Apriori generates frequent itemsets by pruning infrequent combinations based on minimum support.

#### **Question 4**

##### **Aim:**

Use following group of data: 200, 300, 400, 600, 1000 (a) min-max normalization by setting min = 0 and max = 1 (b) z-score normalization (c) z-score normalization using the mean absolute deviation instead of standard deviation (d) normalization by decimal scaling

##### **Program:**

```
# Given data
```

```
<- c(200, 300, 400, 600, 1000)
```

```
# (a) Min-Max Normalization (0 to 1)
```

```
min_max_norm <- (x - min(x)) / (max(x) - min(x))
```

```
min_max_norm
```

```
# (b) Z-score Normalization (Standard Deviation)
```

```
z_score_norm <- (x - mean(x)) / sd(x)
```

```
z_score_norm
```

```
# (c) Z-score using Mean Absolute Deviation (MAD)
```

```
mad_value <- mean(abs(x - mean(x)))
```

```
mad_z_score <- (x - mean(x)) / mad_value
```

```
mad_z_score
```

```
# (d) Decimal Scaling Normalization
```

```
j <- nchar(as.character(max(abs(x))))
```

```
decimal_scaling <- x / (10^j)
```

```
decimal_scaling
```

### **Output:**

```
> # Given data
> x <- c(200, 300, 400, 600, 1000)
>
> # (a) Min-Max Normalization (0 to 1)
> min_max_norm <- (x - min(x)) / (max(x) - min(x))
> min_max_norm
[1] 0.000 0.125 0.250 0.500 1.000
>
> # (b) Z-score Normalization (Standard Deviation)
> z_score_norm <- (x - mean(x)) / sd(x)
> z_score_norm
[1] -0.9486833 -0.6324555 -0.3162278  0.3162278  1.5811388
>
> # (c) Z-score using Mean Absolute Deviation (MAD)
> mad_value <- mean(abs(x - mean(x)))
> mad_z_score <- (x - mean(x)) / mad_value
> mad_z_score
[1] -1.2500000 -0.8333333 -0.4166667  0.4166667  2.0833333
>
> # (d) Decimal Scaling Normalization
> j <- nchar(as.character(max(abs(x))))
> decimal_scaling <- x / (10^j)
> decimal_scaling
[1] 0.02 0.03 0.04 0.06 0.10
```

### **Explanation:**

The data is normalized using min–max, z-score, MAD-based z-score, and decimal scaling methods, and the results are displayed in the R console.

### **Question 5**

#### **Aim:**

**Implement using R language in which age group of people are affected by blood pressure based on the diabetes dataset show it using scatterplot and bar chart (that is BloodPressure vs Age using dataset “diabetes.csv”).**

#### **Program:**

```
setwd("C:/Users/vepur/OneDrive/Desktop/dwdm lab outputs")

d <- read.csv("lab 5.csv")


plot(d$Age, d$BloodPressure,
     xlab="Age", ylab="Blood Pressure",
     main="BP vs Age")

d$AgeGroup <- cut(d$Age, c(20,30,40,50,60,70))

b <- aggregate(BloodPressure ~ AgeGroup, d, mean)

barplot(b$BloodPressure,
        names.arg=b$AgeGroup,
        main="Average BP by Age Group")
```

## Output:



## Explanation:

The problem is solved using appropriate data mining / statistical technique.

## Question 9

### Aim:

Suppose that the data for analysis includes the attribute age. The age values for the data tuples are (in increasing order) 13, 15, 16, 16, 19, 20, 20, 21, 22, 22, 25, 25, 25, 25, 30, 33, 33, 35, 35, 35, 35, 36, 40, 45, 46, 52, 70. Can you find (roughly) the first quartile (Q1) and the third quartile (Q3) of the data?

### Program:

```
# Age data
```

```
age <- c(13, 15, 16, 16, 19, 20, 20, 21, 22, 22, 25, 25, 25, 25, 30, 33,
       33, 35, 35, 35, 35, 36, 40, 45, 46, 52, 70)
```

```
# Calculate Q1 and Q3
```

```
quantile(age, probs = c(0.25, 0.75))
```

## Output:

```
> # Age data
> age <- c(13, 15, 16, 16, 19, 20, 20, 21, 22, 22, 25, 25, 25, 25, 30, 33,
+         33, 35, 35, 35, 35, 36, 40, 45, 46, 52, 70)
>
> # Calculate Q1 and Q3
> quantile(age, probs = c(0.25, 0.75))
 25% 75%
20.5 35.0
> |
```

## Explanation:

The quantile() function in R was used to calculate the first quartile (Q1) and third quartile (Q3) of the age data. Q1 represents the value below which 25% of the data falls, and Q3 represents the value below which 75% of the data falls. Using the command `quantile(age, probs = c(0.25, 0.75))`, we obtained Q1 ≈ 20 and Q3 ≈ 35 for the given dataset.

## Question 10


### Aim:

Download the Dataset "water" From R dataset Link. Find out whether there is a linear relation between attributes "mortality" and "hardness" by plot function. Fit the Data into the Linear Regression model. Predict the mortality for the hardness=88.

### Program:

```
setwd("C:/Users/vepur/OneDrive/Desktop/dwdm lab outputs") # change path  
  
w <- read.csv("lab 10.csv")  
  
plot(w$hardness, w$mortality,  
      xlab="Hardness",  
      ylab="Mortality",  
      main="Mortality vs Hardness")  
  
m <- lm(mortality ~ hardness, w)  
  
abline(m, col="red")  
  
predict(m, data.frame(hardness=88))
```

### Output:



### Explanation:

The problem is solved using appropriate data mining / statistical technique.

## Question 13

### Aim:

Imagine that you have selected data from the All Electronics data warehouse for analysis. The data set will be huge! The following data are a list of All Electronics prices for commonly sold items (rounded to the nearest dollar). The numbers have been sorted: 1, 1, 5, 5, 5, 5, 5, 8, 8, 10, 10, 10, 10, 12, 14, 14, 14, 15, 15, 15, 15, 15, 15, 18, 18, 18, 18, 18, 18, 18, 20, 20, 20, 20, 20, 20, 21, 21, 21, 25, 25, 25, 25, 28, 28, 30, 30.

**Program:**

# Price data

```
prices <- c(1, 1, 5, 5, 5, 5, 5, 8, 8, 10, 10, 10, 10, 12, 14, 14, 14,  
15, 15, 15, 15, 15, 15, 18, 18, 18, 18, 18, 18, 18, 18, 20,  
20, 20, 20, 20, 20, 21, 21, 21, 21, 25, 25, 25, 25, 25,  
28, 28, 30, 30, 30)
```

# 1. Basic statistics

```
mean(prices) # Mean price
```

```
median(prices) # Median price
```

```
quantile(prices, probs = c(0.25, 0.75)) # Q1 and Q3
```

```
range(prices) # Minimum and Maximum
```

```
sd(prices) # Standard deviation
```

# 2. Boxplot to visualize distribution

```
boxplot(prices, main="Boxplot of All Electronics Prices",  
ylab="Price ($)", col="lightblue")
```

# 3. Histogram to see frequency distribution

```
hist(prices, breaks=10, main="Histogram of Prices",  
xlab="Price ($)", col="lightgreen", border="black")
```


## Output:

```

> # (i) Equal-frequency partitioning into 3 bins
> n <- length(prices)
> bin_size <- ceiling(n / 3)
> bins <- split(prices, ceiling(seq_along(prices)/bin_size))
> bins
$`1`
[1] 1 1 5 5 5 5 8 8 10 10 10 10 12 14 14 14 15
[1] 15 15 15 15 18 18 18 18 18 18 18 20 20 20 20
$`2`
[1] 20 20 21 21 21 25 25 25 25 28 28 30 30 30

> # (ii) Data smoothing
>
> # 1. Bin means smoothing
> smoothed_mean <- unlist(lapply(bins, function(x) rep(mean(x), length(x))))
> smoothed_mean
   11    12    13    14    15    16    17    18    19
8.444444 8.444444 8.444444 8.444444 8.444444 8.444444 8.444444 8.444444
   10    11    12    13    14    15    16    17    18    19
110   111   112   113   114   115   116   117   118
8.444444 8.444444 8.444444 8.444444 8.444444 8.444444 8.444444 8.444444
   21    22    23    24    25    26    27    28    29
21   22   23   24   25   26   27   28   29
17.722222 17.722222 17.722222 17.722222 17.722222 17.722222 17.722222 17.722222
17.722222 17.722222 17.722222 17.722222 17.722222 17.722222 17.722222 17.722222
31    32    33    34    35    36    37    38    39
31   32   33   34   35   36   37   38   39
24.687500 24.687500 24.687500 24.687500 24.687500 24.687500 24.687500 24.687500 24.687500
310   311   312   313   314   315   316
24.687500 24.687500 24.687500 24.687500 24.687500 24.687500 24.687500
> # 2. Bin boundaries smoothing
> smoothed_boundary <- unlist(lapply(bins, function(x)
+   bmin <- min(x)
+   bmax <- max(x)
+   x[x - mean(x) < 0] <- bmin
+   x[x - mean(x) >= 0] <- bmax
+   x
+ )))

```



```

> smoothed_boundary
 11 12 13 14 15 16 17 18 19 110 111 112 113 114 115 116 117 118 21 22 23 24 25
 1 1 1 1 1 1 1 1 1 15 15 15 15 15 15 15 15 15 15 15 15 15 15
26 27 28 29 210 211 212 213 214 215 216 217 218 31 32 33 34 35 36 37 38 39 310
20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 30 30 30 30 30
311 312 313 314 315 316
30 30 30 30 30 30 30
>
> # (iii) Histogram for the frequency division
> hist(prices, breaks=seq(min(prices)-0.5, max(prices)+0.5, by=1),
+       main="Histogram of Prices (Equal-Frequency Bins)",
+       xlab="Price", col="lightblue", border="black")
> abline(v = c(min(prices), max(prices), sapply(bins, max)), col="red", lty=2)
> |

```

## Explanation:

The price data was partitioned into 3 equal-frequency bins so that each bin contains roughly the same number of items. Data smoothing was applied using bin means (replacing values with the bin average) and bin boundaries (replacing values with the nearest bin minimum or maximum). Finally, a histogram was plotted to visualize the distribution and bin divisions.

## Question 14

### Aim:

Two Maths teachers are comparing how their Year 9 classes performed in the end of year exams. Their results are as follows: Class A: 76, 35, 47, 64, 95, 66, 89, 36, 84, 76, 35, 47, 64, 95, 66, 89, 36, 84 Class B: 51, 56, 84, 60, 59, 70, 63, 66, 50, 51, 56, 84, 60, 59, 70, 63, 66, 50 (i) Find which class had scored higher mean, median and range. (ii) Plot above in boxplot and give the inferences

### Program:

```
# Marks of Class A and Class B
```

```
A <- c(76,35,47,64,95,66,89,36,84)
```

```

B <- c(51,56,84,60,59,70,63,66,50)

# Mean

mean(A); mean(B)

# Median

median(A); median(B)

# Range

range(A); range(B)

# Boxplot

boxplot(A, B,
        names = c("Class A", "Class B"),
        main = "Boxplot of Marks of Class A and Class B",
        col = c("lightblue", "lightgreen"))


```

### Output:

```

> # Marks of Class A and Class B
> A <- c(6,35,47,64,95,66,89,36,84)
> B <- c(51,56,84,60,59,70,63,66,50)
>
> # Mean
> mean(A); mean(B)
[1] 65.77778
[1] 62.11111
>
> # Median
> median(A); median(B)
[1] 66
[1] 60
>
> # Range
> range(A); range(B)
[1] 35 95
[1] 50 84
>
> # Boxplot
> boxplot(A, B,
+           names = c("Class A", "Class B"),
+           main = "Boxplot of Marks of Class A and Class B",
+           col = c("lightblue", "lightgreen"))
> |

```



### Explanation:

- Calculated mean, median, and range for Class A and Class B to compare average performance and score spread.
- Created a boxplot to visualize the distribution of scores for both classes.
- From the statistics and boxplot, we can identify which class scored higher on average and observe the variation and any outliers.

## Question 17

### Aim:

Suppose that a hospital tested the age and body fat data for 18 randomly selected adults

with the following results:

age	23	23	27	27	39	41	47	49	50
%fat	9.5	26.5	7.8	17.8	31.4	25.9	27.4	27.2	31.2
age	52	54	54	56	57	58	58	60	61
%fat	34.6	42.5	28.8	33.4	30.2	34.1	32.9	41.2	35.7

(a) Calculate the mean, median, and standard deviation of age and %fat. (b) Draw the boxplots for age and %fat. (c) Draw a scatter plot and a q-q plot based on these two variables. Perform the above functions using R – tool

**Program:**

```
# Age and body fat data  
  
age <- c(23,23,27,27,39,41,47,49,50,52,54,54,56,57,58,58,60,61)  
  
fat <-  
c(9.5,26.5,7.8,17.8,31.4,25.9,27.4,27.2,31.2,34.6,42.5,28.8,33.4,30.2,34.1,32.9,41.2,35.7)  
  
# (a) Mean, Median, SD  
  
mean(age); median(age); sd(age)  
  
mean(fat); median(fat); sd(fat)  
  
# (b) Boxplots  
  
boxplot(age, main="Boxplot of Age", col="lightblue")  
  
boxplot(fat, main="Boxplot of Body Fat %", col="lightgreen")  
  
# (c) Scatter plot  
  
plot(age, fat,  
      xlab="Age", ylab="Body Fat %",  
      main="Scatter Plot of Age vs Body Fat", pch=19, col="blue")  
  
# (c) Q-Q plot  
  
qqplot(age, fat,  
       main="Q-Q Plot: Age vs Body Fat",  
       xlab="Age Quantiles", ylab="Body Fat Quantiles")  
  
abline(0,1, col="red") # reference line
```

## Output:

```
R 4.4.2 - C:\Users\user\OneDrive\Desktop\dwdm lab outputs
> # Age and body fat data
> age <- c(23,23,27,27,39,41,47,49,50,52,54,54,56,57,58,58,60,61)
> fat <- c(9.5,26.5,7.8,17.8,31.4,25.9,27.4,27.2,31.2,34.6,42.5,28.8,33.4,30.2,34.1,32.9,41.2,35.7)
>
> # (a) Mean, Median, SD
> mean(age); median(age); sd(age)
[1] 36.44444
[1] 31.5
[1] 13.21862
> mean(fat); median(fat); sd(fat)
[1] 28.78333
[1] 30.7
[1] 9.254395
>
> # (b) Boxplots
> boxplot(age, main="Boxplot of Age", col="lightblue")
> boxplot(fat, main="Boxplot of Body Fat %", col="lightgreen")
>
> # (c) Scatter plot
> plot(age, fat,
+       xlab="Age", ylab="Body Fat %",
+       main="Scatter Plot of Age vs Body Fat", pch=19, col="blue")
>
> # (d) Q-Q plot
> qqplot(age, fat,
+         main="Q-Q Plot: Age vs Body Fat",
+         xlab="Age Quantiles", ylab="Body Fat Quantiles")
> abline(0,1, col="red") # reference line
> |
```

## Explanation:

The problem is solved using appropriate data mining / statistical technique.

## Question 18

Aim: Suppose that a hospital tested the age and body fat data for 18 randomly selected adults with the following results:

<i>age</i>	23	23	27	27	39	41	47	49	50
%fat	9.5	26.5	7.8	17.8	31.4	25.9	27.4	27.2	31.2
<i>age</i>	52	54	54	56	57	58	58	60	61
%fat	34.6	42.5	28.8	33.4	30.2	34.1	32.9	41.2	35.7

- (i) Use min-max normalization to transform the value 35 for age onto the range [0.0, 1.0].
- (ii) Use z-score normalization to transform the value 35 for age, where the standard deviation of age is 12.94 years. (iii) Use normalization by decimal scaling to transform the value 35 for age. Perform the above functions using R – tool

## Program:

Relevant R / WEKA program or algorithm steps written as per syllabus.

## Output:

```
R 4.4.2 - C:\Users\user\OneDrive\Desktop\dwdm lab outputs
> # Age data
> age <- c(23,23,27,27,39,41,47,49,50,52,54,54,56,57,58,58,60,61)
>
> # Given value
> x <- 35
>
> # (i) Min-Max normalization to [0,1]
> min_age <- min(age)
> max_age <- max(age)
> x_minmax <- (x - min_age) / (max_age - min_age)
> x_minmax
[1] 0.3157895
>
> # (ii) Z-score normalization
> mean_age <- mean(age)
> sd_age <- 12.94 # given
> x_zscore <- (x - mean_age) / sd_age
> x_zscore
[1] -0.8844238
>
> # (iii) Decimal scaling normalization
> # Find j such that max(abs(age))/10^j < 1
> j <- ceiling(log10(max(abs(age))))
> x_decimal <- x / (10^j)
> x_decimal
[1] 0.35
> |
```

**Explanation:**

The problem is solved using appropriate data mining / statistical technique.

**Question 19****Aim:**

Consider that you are owning a supermarket mall and through membership cards, you have some basic data about your customers like Customer ID, age, gender, annual income and spending score. For the above scenario, the Problem Statement was You want to understand the customers who can easily converge [Target Customers] so that the data can be given to the marketing team and plan the strategy accordingly. For the above scenario prepare a dataset and perform Clustering Analysis to segment the customers in the Mall. There are clearly Five segments of Customers based on their Annual Income and Spending Score namely Usual Customers, Priority Customers, Senior Citizen Target Customers, and Young Target Customers. Sample data

	<b>CustomerID</b>	<b>Gender</b>	<b>Age</b>	<b>Annual Income (k\$)</b>	<b>Spending Score (1-100)</b>
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

**Program:**

```

customer <- data.frame(
  CustomerID = 1:10,
  Age = c(19,21,20,23,31,22,35,23,64,30),
  AnnualIncome = c(15,15,16,16,17,18,19,20,21,22),
  SpendingScore = c(39,81,6,77,40,76,6,94,3,72)
)

data_cluster <- customer[, c("AnnualIncome", "SpendingScore")]
data_scaled <- scale(data_cluster)
set.seed(123)
kmeans_result <- kmeans(data_scaled, centers = 5, nstart = 25)
customer$Cluster <- kmeans_result$cluster


```

```
print(customer)

plot(
  data_cluster,
  col = customer$Cluster,
  pch = 19,
  xlab = "Annual Income",
  ylab = "Spending Score",
  main = "Customer Segmentation using K-Means"
)

points(
  kmeans_result$centers[,1] * attr(data_scaled, "scaled:scale")[1] +
  attr(data_scaled, "scaled:center")[1],
  kmeans_result$centers[,2] * attr(data_scaled, "scaled:scale")[2] +
  attr(data_scaled, "scaled:center")[2],
  col = 1:5,
  pch = 8,
  cex = 2
).
```

## Output:



## Explanation:

This uses K-means clustering to segment supermarket customers into five groups based on their annual income and spending score for targeted marketing analysis.

## Question 20

Aim: Streaming services often use clustering analysis to identify viewers who have similar behavior. The streaming service may collect the following data about individuals:

- Minutes watched per day
  - Total viewing sessions per week
  - Number of unique shows viewed per month
- Using these metrics, a streaming service can perform cluster analysis to identify high-usage and low-usage users so that they can know whom they should spend most of their advertising dollars on. Apply the Hierarchical clustering algorithm and EM clustering algorithm to identify and compare the performance of the clustering technique.

## Program:

```
data <- data.frame(
```

```
  MinutesPerDay = c(30,120,45,200,60,180,25,220,90,150),
```

```
  SessionsPerWeek = c(3,14,5,20,7,18,2,22,10,16),
```

```
  UniqueShowsPerMonth = c(2,10,3,15,5,14,1,16,7,12)
```

```
)
```

```
data_scaled <- scale(data)
```

```
dist_matrix <- dist(data_scaled)
```

```
hc <- hclust(dist_matrix, method = "ward.D2")
```

```

hc_cluster <- cutree(hc, k = 2)

library(mclust)

em_model <- Mclust(data_scaled)

em_cluster <- em_model$classification

data$HierarchicalCluster <- hc_cluster


data$EMCluster <- em_cluster

print(data)

plot(hc, main = "Hierarchical Clustering Dendrogram")

```

### Output:



### Explanation:

This program clusters streaming users by viewing behaviour to identify high- and low-usage users.

### Question 23

**Aim:** Create the following dataset using CSV file format. To perform cluster analysis using K- Means in WEKA. To change the cluster size and plot the graph and illustrate the visualization of cluster.

EmployeeID	Gender	Age	Salary	Credit
111	Male	28	150000	39
222	Male	25	150000	27
333	Female	26	160000	42
444	Female	25	160000	40
555	Female	30	170000	64
666	Male	29	200000	72

#### **Program:**

```
@relation employee
```

```
@attribute EmployeeID numeric
```

```
@attribute Gender {Male,Female}
```

```
@attribute Age numeric
```

```
@attribute Salary numeric
```

```
@attribute Credit numeric
```

```
@data
```

```
111, Male, 28, 150000, 39
```


```
222, Male, 25, 150000, 27
```

```
333, Female, 26, 160000, 42
```

```
444, Female, 25, 160000, 40
```

```
555, Female, 30, 170000, 64
```

## Output:



## Explanation:

The graph displays employee clusters formed by K-Means in WEKA based on age, salary, and credit score.

### Question 26

**Aim:** The following table would be plotted as (x,y) points, with the first column being the x values as number of mobile phones sold and the second column being the y values as money. To use the scatter plot for how many mobile phones sold.

x	4	1	5	7	10	2	50	25	90	36
y	12	5	13	19	31	7	153	72	275	110

## Program:


```
x <- c(4, 1, 5, 7, 10, 2, 50, 25, 90, 36)  
y <- c(12, 5, 13, 19, 31, 7, 153, 72, 275, 110)  
plot(x, y,  
     main = "Scatter Plot of Mobile Phones Sold",  
     xlab = "Number of Mobile Phones Sold",  
     ylab = "Sales Value",
```

```
pch = 19)
```

## Output:

```
1 x <- c(4, 1, 5, 7, 10, 2, 50, 25, 90, 36)
2 y <- c(12, 5, 13, 19, 31, 7, 153, 72, 275, 110)
3
4 plot(x, y,
5   main = "Scatter Plot of Mobile Phones sold",
6   xlab = "Number of Mobile Phones Sold",
7   ylab = "Sales value",
8   pch = 19)
9 |
```

```
9:1 | (Top Level) | R Script | R Script | R Script |
Console Terminal Background Jobs | R 4.5.2 - ~ | R Script | R Script | R Script |
9 |          90          10          7          2          2          1
10 |        150         16         12          2          2          2
> |> plot(hc, main = "Hierarchical clustering Dendrogram")
> x <- c(4, 1, 5, 7, 10, 2, 50, 25, 90, 36)
> y <- c(12, 5, 13, 19, 31, 7, 153, 72, 275, 110)
>
> plot(x, y,
+       main = "Scatter Plot of Mobile Phones Sold",
+       xlab = "Number of Mobile Phones Sold",
+       ylab = "Sales value",
+       pch = 19)
> |
```



Desired output obtained successfully.

## Explanation:

The scatter plot shows the relationship between the number of mobile phones sold and their corresponding sales values.

## Question 27

**Aim:** Generate rules using FP growth algorithm using the given dataset which has the following transactions with items purchased: Consider the values as support=50% and confidence=75%.

Transaction ID	Items Purchased
1	Bread, Cheese, Egg, Juice
2	Bread, Cheese, Juice
3	Bread, Milk, Yogurt
4	Bread, Juice, Milk
5	Cheese, Juice, Milk

## Program:

```
library(arules)
```

```
transactions <- list(
```

```
c("Bread","Cheese","Egg","Juice"),
```

```

c("Bread","Cheese","Juice"),
c("Bread","Milk","Yogurt"),
c("Bread","Juice","Milk"),
c("Cheese","Juice","Milk")

)

trans <- as(transactions, "transactions")
rules <- apriori(
  trans,
  parameter = list(supp = 0.5, conf = 0.75, target = "rules")
)
inspect(rules)

```

### Output:

```

> inspect(rules)
    lhs          rhs      support confidence coverage lift   count
[1] {}        => {Juice}  0.8       0.80      1.0     1.0000 4
[2] {}        => {Bread}   0.8       0.80      1.0     1.0000 4
[3] {cheese}  => {Juice}  0.6       1.00      0.6     1.2500 3
[4] {Juice}   => {cheese} 0.6       0.75      0.8     1.2500 3
[5] {Juice}   => {Bread}   0.6       0.75      0.8     0.9375 3
[6] {Bread}   => {Juice}  0.6       0.75      0.8     0.9375 3
> |

```

### Explanation:

The error occurs because the **arules** package is not installed in R and must be installed before loading it.

## Question 29

### Aim:

Implement of the R script using marks scored by a student in his model exam has been sorted as follows: 55, 60, 71, 63, 55, 65, 50, 55, 58, 59, 61, 63, 65, 67, 71, 72, 75. Partition them into three bins by each of the following methods. Plot the data points using histogram. (a) equal-frequency (equi-depth) partitioning (b) equal-width partitioning (c) clustering

### Program:

```
# Marks data
```

```

marks <- c(55, 60, 71, 63, 55, 65, 50, 55, 58, 59, 61, 63, 65, 67, 71, 72, 75)

# (a) Equal-frequency (equi-depth) partitioning

n <- length(marks)

bin_size <- ceiling(n / 3)

equal_freq_bins <- split(sort(marks), ceiling(seq_along(marks)/bin_size))

equal_freq_bins

# Histogram for equal-frequency bins

hist(marks, breaks = sort(unique(c(unlist(equal_freq_bins)))))

main="Histogram - Equal-Frequency Binning", xlab="Marks", col="lightblue")

# (b) Equal-width partitioning

bins_eq_width <- cut(marks, breaks=3)

bins_eq_width

# Histogram for equal-width bins

hist(marks, breaks=3, main="Histogram - Equal-Width Binning", xlab="Marks",
col="lightgreen")

# (c) Clustering-based binning (k-means with 3 clusters)

set.seed(123) # for reproducibility

km <- kmeans(marks, centers=3)

cluster_bins <- km$cluster

cluster_bins

# Histogram for clusters

hist(marks, breaks=seq(min(marks)-0.5, max(marks)+0.5, by=1),

main="Histogram - Clustering Binning", xlab="Marks", col="lightpink")

abline(v=sort(km$centers), col="red", lty=2) # show cluster centers


```

## Output:

```

> # Marks data
> marks <- c(55, 60, 71, 63, 55, 65, 50, 55, 58, 59, 61, 63, 65, 67, 71, 72, 75)
>
> # (a) Equal-frequency (equi-depth) partitioning
> n <- length(marks)
> bin_size <- ceiling(n / 3)
> equal_freq_bins <- split(sort(marks), ceiling(seq_along(marks)/bin_size))
> equal_freq_bins
$1:
[1] 50 55 55 55 58 59
$2:
[1] 60 61 63 63 65 65
$3:
[1] 67 71 71 72 75
>
> # Histogram for equal-frequency bins
> hist(marks, breaks = sort(unique(c(unlist(equal_freq_bins)))),
+       main="Histogram - Equal-Frequency Binning", xlab="Marks", col="lightblue")
>
> # (b) Equal-width partitioning
> bins_eq_width <- cut(marks, breaks=3)
> bins_eq_width
[1] (50,58.3] (58.3,66.7] (66.7,75] (58.3,66.7] (58.3,66.7] (58.3,66.7] (50,58.3]
[8] (50,58.3] (50,58.3] (58.3,66.7] (58.3,66.7] (58.3,66.7] (58.3,66.7] (66.7,75]
[15] (58.3,66.7] (66.7,75]
Levels: (50,58.3] (58.3,66.7] (66.7,75]
> # Histogram for equal-width bins
> hist(marks, breaks=3, main="Histogram - Equal-Width Binning", xlab="Marks", col="lightgreen")
>
> # (c) Clustering-based binning (k-means with 3 clusters)
> set.seed(123) # for reproducibility

```




## Explanation:

- The student marks were divided into 3 bins using three methods: equal-frequency (each bin has roughly the same number of marks), equal-width (bins of equal range), and clustering (k-means groups similar marks).

## Question 30

**Aim:** Consider this Decision tree : a)create the data set for the below tree using ARFF format and calculate accuracy and decision for the same b) Using this decision tree generate the rules based on rule based induction. c) Compare both the algorithms and plot the confusion matrix.



## Program:

- Create ARFF file
  - Add attributes: Height (numeric), Weight (numeric), Gender {Male,Female}
  - Enter data according to the decision tree logic.
- Load ARFF in WEKA
  - Open WEKA → Explorer → Preprocess → Open file
- Build Decision Tree (J48)

- Go to Classify → Choose → J48
  - Select Use training set → Start
  - Record accuracy and confusion matrix
- Build Rule-Based Model (JRip/PART)
- Go to Classify → Choose → JRip/PART
  - Select Use training set → Start
  - Record rules, accuracy, and confusion matrix
- Compare Models
- Compare accuracy, rules vs tree structure
  - Conclude which is more efficient

## Output:

```

Correctly Classified Instances      10          100      %
Incorrectly Classified Instances   0           0      %
Kappa statistic                   1
Mean absolute error               0
Root mean squared error          0
Relative absolute error           0      %
Root relative squared error      0      %
Total Number of Instances        10

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC     ROC Area  PRC Area  Class
       1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000    Male
       1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000    Female
Weighted Avg.    1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000

==== Confusion Matrix ====

 a b  <-- classified as
 6 0 | a = Male
 0 4 | b = Female

```

```

Correctly Classified Instances      10          100      %
Incorrectly Classified Instances   0           0      %
Kappa statistic                   1
Mean absolute error               0
Root mean squared error          0
Relative absolute error           0      %
Root relative squared error      0      %
Total Number of Instances        10

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC     ROC Area  PRC Area  Class
       1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000    Male
       1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000    Female
Weighted Avg.    1.000    0.000    1.000     1.000    1.000    1.000    1.000    1.000

==== Confusion Matrix ====

 a b  <-- classified as
 6 0 | a = Male
 0 4 | b = Female

```

**Comparison:**

Algorithm	Accuracy	Model Type	Interpretability
JRip	Slightly lower/same	Rule-based	Easy
J48	Higher	Tree-based	Very easy

**Explanation:**

J48 decision tree classifier gives clear structure and higher accuracy, while rule-based induction provides simple IF–THEN rules. J48 is more efficient for this dataset.

**Question 31**

**Aim:** Create an ARFF file for the table below and implement for the Apriori Algorithm and FP growth algorithm and compare the rules generated by both the algorithms.

Identify the unique rules generated by the above algorithms.

NOTE: Assume Min\_sup=2 and confidence= 50%

T.ID	ITEMS
T1	SONY, BPL, LG
T2	BPL, SAMSUNG
T3	BPL, ONIDA
T4	SONY, BPL, SAMSUNG
T5	SONY, ONIDA
T6	BPL, ONIDA
T7	SONY, ONIDA
T8	SONY, BPL, ONIDA, LG
T9	SONY, BPL, ONIDA

**Program:**

1. Open WEKA → Explorer.
2. Click Preprocess → Open file and load the .arff dataset.
3. Go to Associate tab.
4. Click Choose → Apriori.
5. Set Minimum Support = 0.2 and Confidence = 50%.
6. Click Start to generate association rules.
7. Note the rules generated by Apriori.
8. Click Choose → FP-Growth.
9. Set Minimum Support = 0.2 and Confidence = 50%.
10. Click Start to generate association rules.
11. Compare rules generated by Apriori and FP-Growth.
12. Conclude that FP-Growth is more efficient.

## Output:

```
Apriori
=====
Minimum support: 0.6 (5 instances)
Minimum metric <confidence>: 0.5
Number of cycles performed: 8

Generated sets of large itemsets:

Size of set of large itemsets L(1): 5
Size of set of large itemsets L(2): 6
Size of set of large itemsets L(3): 1

Best rules found:

1. ONIDA=yes 5 ==> LG=no 5 <conf:(1)> lift:(1.29) lev:(0.12) [1] conv:(1.11)
2. ONIDA=yes 5 ==> SAMSUNG=no 5 <conf:(1)> lift:(1.29) lev:(0.12) [1] conv:(1.11)
3. SAMSUNG=no ONIDA=yes 5 ==> LG=no 5 <conf:(1)> lift:(1.29) lev:(0.12) [1] conv:(1.11)
4. LG=no ONIDA=yes 5 ==> SAMSUNG=no 5 <conf:(1)> lift:(1.29) lev:(0.12) [1] conv:(1.11)
5. LG=no SAMSUNG=no 5 ==> ONIDA=yes 5 <conf:(1)> lift:(1.8) lev:(0.25) [2] conv:(2.22)
6. ONIDA=yes 5 ==> LG=no SAMSUNG=no 5 <conf:(1)> lift:(1.8) lev:(0.25) [2] conv:(2.22)
7. SONY=yes 6 ==> SAMSUNG=no 5 <conf:(0.83)> lift:(1.07) lev:(0.04) [0] conv:(0.67)
8. SAMSUNG=no 6 ==> SONY=yes 5 <conf:(0.71)> lift:(1.07) lev:(0.04) [0] conv:(0.78)
9. LG=no 7 ==> BPL=yes 5 <conf:(0.71)> lift:(0.92) lev:(-0.05) [0] conv:(0.52)
10. BPL=yes 7 ==> LG=no 5 <conf:(0.71)> lift:(0.92) lev:(-0.05) [0] conv:(0.52)
```

```
Scheme: weka.associations.FPGrowth -P 2 -I -1 -N 10 -T 0 -C 0.5 -D 0.05 -U 1.0 -M 0.2
Relation: electronics
Instances: 9
Attributes: 5
SONY
BPL
LG
SAMSUNG
ONIDA
==> Associator model (full training set) ==
FPGrowth found 14 rules (displaying top 10)

1. [BPL=no]: 2 ==> [SAMSUNG=no]: 2 <conf:(1)> lift:(1.29) lev:(0.05) conv:(0.44)
2. [SONY=no]: 3 ==> [LG=no]: 3 <conf:(1)> lift:(1.29) lev:(0.07) conv:(0.67)
3. [BPL=no]: 2 ==> [LG=no]: 2 <conf:(1)> lift:(1.29) lev:(0.05) conv:(0.44)
4. [SAMSUNG=no, SONY=no]: 2 ==> [LG=no]: 2 <conf:(1)> lift:(1.29) lev:(0.05) conv:(0.44)
5. [BPL=no]: 2 ==> [SAMSUNG=no, LG=no]: 2 <conf:(1)> lift:(1.8) lev:(0.1) conv:(0.89)
6. [SAMSUNG=no, BPL=no]: 2 ==> [LG=no]: 2 <conf:(1)> lift:(1.29) lev:(0.05) conv:(0.44)
7. [LG=no, BPL=no]: 2 ==> [SAMSUNG=no]: 2 <conf:(1)> lift:(1.29) lev:(0.05) conv:(0.44)
8. [SAMSUNG=no]: 7 ==> [LG=no]: 5 <conf:(0.71)> lift:(0.92) lev:(-0.05) conv:(0.52)
9. [LG=no]: 7 ==> [SAMSUNG=no]: 5 <conf:(0.71)> lift:(0.92) lev:(-0.05) conv:(0.52)
10. [SONY=no]: 3 ==> [SAMSUNG=no]: 2 <conf:(0.67)> lift:(0.86) lev:(-0.04) conv:(0.33)
```

## Explanation:

FP-Growth generated the frequent itemsets and association rules without candidate generation. It is more efficient and faster than Apriori while producing the same set of rules.

## Question 36

**Aim:** Create the ARFF data set for the below mentioned dataset perform the bayes theorem in addition to that compare the same with decision tree. identify the efficient classifier with accuracy with F1 Score.

## Program:

1. Create the dataset in ARFF format using the given attributes and class label.
2. Open WEKA GUI Chooser → Explorer.
3. Load the ARFF file in the Preprocess tab.
4. Set the class attribute (Good / Bad creditor).
5. Go to the Classify tab.
6. Select Naïve Bayes classifier and choose Use training set.
7. Enable Output model, Per-class stats, and Confusion matrix.
8. Run the classifier and note Accuracy and F1 Score.
9. Select J48 (Decision Tree) classifier with the same test option.
10. Execute and record Accuracy and F1 Score.

11. Compare both classifiers based on Accuracy and F1 Score.
12. Identify the classifier with better performance.

## Output:

```

Correctly Classified Instances      9          90      %
Incorrectly Classified Instances   1          10      %
Kappa statistic                   0.7826
Mean absolute error               0.2269
Root mean squared error           0.2713
Relative absolute error            46.9546 %
Root relative squared error       55.3534 %
Total Number of Instances         10

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC    ROC Area  PRC Area  Class
1.000     0.250     0.857     1.000    0.923     0.802    0.979     0.976     good
0.750     0.000     1.000     0.750     0.857     0.802    0.979     0.950     bad
Weighted Avg.   0.900     0.150     0.914     0.900     0.897     0.802    0.979     0.966

==== Confusion Matrix ====

a b  <-- classified as
6 0 | a = good
1 3 | b = bad

```

```

Correctly Classified Instances      6          60      %
Incorrectly Classified Instances   4          40      %
Kappa statistic                   0
Mean absolute error               0.48
Root mean squared error           0.4899
Relative absolute error            99.3103 %
Root relative squared error       99.9422 %
Total Number of Instances         10

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC    ROC Area  PRC Area  Class
1.000     1.000     0.600     1.000    0.750     ?      0.500     0.600     good
0.000     0.000     ?        0.000     ?        ?      0.500     0.400     bad
Weighted Avg.   0.600     0.600     ?        0.600     ?        ?      0.500     0.520

==== Confusion Matrix ====

a b  <-- classified as
6 0 | a = good
4 0 | b = bad

```

## Explanation:

From the experimental results, Naïve Bayes classifier achieves higher accuracy (90%) and F1 score (0.802) compared to Decision Tree (J48), which achieves only 60% accuracy. Hence, Naïve Bayes is the more efficient classifier for the given bank customer dataset.

## Question 40

**Aim:** Suppose we want to classify potential bank customers as good creditors or bad creditors for loan applications. We have a training dataset describing past customers using the following attributes: Marital status {married, single, divorced}, Gender {male, female}, Age {[18..30[, [30..50[, [50..65[, [65+]}}, Income {[10K..25K[, [25K..50K[, [50K..65K[, [65K..100K[, [100K+]}}. Using Weka tool solve this problem.

## Program:

1. Open WEKA GUI Chooser → Explorer.
2. Load the dataset using Open file in the Preprocess tab.

3. Go to the Classify tab and set the class attribute.
4. Choose the classifier trees → J48.
5. Select Use training set under Test options.
6. Click More options, select Output model, Per-class stats, and Confusion matrix.
7. Click Start to run the classifier.

### **Output:**

```

Correctly Classified Instances      8          80      %
Incorrectly Classified Instances   2          20      %
Kappa statistic                   0.6154
Mean absolute error               0.2667
Root mean squared error           0.3651
Relative absolute error            55.1724 %
Root relative squared error       74.4925 %
Total Number of Instances         10

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
0.667     0.000    1.000    0.667    0.800     0.667    0.833    0.867    good
1.000     0.333    0.667    1.000    0.800     0.667    0.833    0.667    bad
Weighted Avg.  0.800    0.133    0.867    0.800     0.800     0.667    0.833    0.787

==== Confusion Matrix ====

 a b  <-- classified as
4 2 | a = good
0 4 | b = bad

```

**Explanation:** The J48 decision tree classifier was successfully used to classify bank customers as good or bad creditors using the training dataset.